

PCA

June 6, 2026

1 PCA First

```
[24]: import pandas as pd
      from sklearn.decomposition import PCA
      from sklearn.cluster import KMeans
      import matplotlib.pyplot as plt
      from sklearn.cluster import DBSCAN
      import seaborn as sns
      from sklearn.metrics import silhouette_score
```

1.1 Loading the data

```
[25]: # Load data and extract features
      df = pd.read_csv("../data/Qualifying_scaled_for_clustering.csv")
      identifiers = ['track', 'global_corner_id', 'driver_number']
      features = df.drop(columns=identifiers)
```

1.2 Applying PCA

```
[26]: # Apply PCA retaining 95% of variance
      pca = PCA(n_components=0.95)
      pca_data = pca.fit_transform(features)

      # Print feature reduction statistics
      orig_count = features.shape[1]
      reduced_count = pca_data.shape[1]
      dropped_count = orig_count - reduced_count

      print(f"Original number of features: {orig_count}")
      print(f"Features after PCA (components): {reduced_count}")
      print(f"Number of features reduced: {dropped_count}")
```

```
Original number of features: 12
Features after PCA (components): 9
Number of features reduced: 3
```

We use PCA to simplify our data by grouping related features into independent components. This fixes the problem of having too many variables, removes statistical noise, and helps the clustering

algorithm run faster and more accurately.

1.3 Finding the ideal number of clusters



```
[27]: wcss = []
silhouette_scores = []
k_values = range(1, 11)

for k in k_values:
    kmeans = KMeans(n_clusters=k, random_state=42)
    cluster_labels = kmeans.fit_predict(pca_data)
    wcss.append(kmeans.inertia_)

    # Silhouette score is only defined for 2 <= k <= n_samples - 1
    if k > 1:
        score = silhouette_score(pca_data, cluster_labels)
        silhouette_scores.append(score)

# Plot WCSS (Elbow Method)
plt.figure(figsize=(10, 6))
plt.plot(k_values, wcss, marker='o', linestyle='--')
plt.title('Elbow Method For Optimal k')
plt.xlabel('Number of clusters (k)')
plt.ylabel('WCSS')
plt.grid(True)
plt.show()

# Plot Silhouette Scores
plt.figure(figsize=(10, 6))
k_values_sil = range(2, 11)
plt.plot(k_values_sil, silhouette_scores, marker='s', linestyle='-', color='green')
plt.title('Silhouette Score For Optimal k')
plt.xlabel('Number of clusters (k)')
plt.ylabel('Silhouette Score')
plt.grid(True)
plt.show()
```



The Silhouette score peaks at $k=2$ (just above 0.2), suggesting it as the mathematical optimum. However, considering the inflection point from the Elbow method, we chose to proceed with $k=3$

to achieve a more practical and meaningful division of real-world driving styles. The relatively low overall Silhouette score indicates overlapping clusters. This is an expected result, as continuous F1 telemetry data reflects fluid driving behaviors without absolute, sharp boundaries.

1.4 K-Means

```
[28]: # Run K-Means and save labels to the original dataframe
kmeans = KMeans(n_clusters=3, random_state=42)
df['kmeans_cluster'] = kmeans.fit_predict(pca_data)
```

1.5 DBSCAN

```
[29]: # Run DBSCAN and save labels to the original dataframe
dbscan = DBSCAN(eps=0.7, min_samples=5)
df['dbscan_cluster'] = dbscan.fit_predict(pca_data)
```



1.6 Visualization - HEATMAP

```
[30]: # Calculate feature means for K-Means
kmeans_means = df.groupby('kmeans_cluster')[features.columns].mean()

# Plot K-Means Heatmap
plt.figure(figsize=(14, 6))
sns.heatmap(kmeans_means.T, annot=True, cmap='coolwarm', fmt=".2f", center=0)
plt.title('K-Means Cluster Profiles')
plt.xlabel('Cluster ID')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()

# Calculate feature means for DBSCAN
dbscan_means = df.groupby('dbscan_cluster')[features.columns].mean()
dbscan_means.index = [f'Noise (-1)' if c == -1 else f'Cluster {c}' for c in
    ↪ dbscan_means.index]

# Plot DBSCAN Heatmap
plt.figure(figsize=(14, 6))
sns.heatmap(dbscan_means.T, annot=True, cmap='coolwarm', fmt=".2f", center=0)
plt.title('DBSCAN Cluster Profiles')
plt.xlabel('Cluster ID')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()
```




We can see that k-means provided clearer and better results, for example, cluster 0 resembles drivers who drive slowly and carefully, with lower speeds overall.

1.7 Visualization - BARPLOT

```
[31]: # Plot K-Means Bar Chart
kmeans_means.T.plot(kind='bar', figsize=(15, 6), colormap='viridis')
plt.title('K-Means: Feature Means per Cluster')
plt.xlabel('Feature')
plt.ylabel('Average Value (Z-Score)')
plt.legend(title='Cluster', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
```

```
plt.show()

# Plot DBSCAN Bar Chart
dbscan_means.T.plot(kind='bar', figsize=(15, 6), colormap='plasma')
plt.title('DBSCAN: Feature Means per Cluster')
plt.xlabel('Feature')
plt.ylabel('Average Value (Z-Score)')
plt.legend(title='Cluster', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```



We can use the bar plot of the K-Means to easily understand the characteristics of the different clusters:


Cluster 0: Low speeds, small braking

Cluster 1: High speed, high braking, aggressive driving

Cluster 2: Low entry speeds, almost no breaking, smooth driving

```
[32]: # 1. Macro: Driver vs. K-Means Cluster distribution
driver_cluster_dist = pd.crosstab(df['driver_number'], df['kmeans_cluster'])

plt.figure(figsize=(12, 6))
sns.heatmap(driver_cluster_dist, annot=True, fmt="d", cmap="Blues")
plt.title('Corner Counts: Driver vs. K-Means Cluster')
plt.xlabel('K-Means Cluster Of Corners')
plt.ylabel('Driver Number')
plt.show()
```



This visualization reveals individual driving habits by highlighting the most common cornering style assigned to each driver.

1.8 Final Clustering

```
[33]: driver_profiles = pd.crosstab(df['driver_number'], df['kmeans_cluster'],
    ↪normalize='index')
driver_profiles.columns = [f'Corner Style {c}' for c in driver_profiles.columns]

kmeans_drivers = KMeans(n_clusters=3, random_state=42)
driver_profiles['Driver_Style_Cluster'] = kmeans_drivers.
    ↪fit_predict(driver_profiles)

driver_profiles_sorted = driver_profiles.sort_values('Driver_Style_Cluster')
```

```

cluster_labels = driver_profiles_sorted.pop('Driver_Style_Cluster')


plt.figure(figsize=(12, 8))
sns.heatmap(driver_profiles_sorted, annot=True, fmt=".1%", cmap="YlGnBu",
            vmin=0, vmax=1)

y_labels = [f"Driver {idx} (Cluster {c})" for idx, c in
            zip(driver_profiles_sorted.index, cluster_labels)]
plt.yticks(ticks=[i + 0.5 for i in range(len(y_labels))], labels=y_labels,
            rotation=0)

plt.title('Driver Profiles: Corner Style Distribution')
plt.ylabel('Driver Number & Style Cluster')
plt.xlabel('Corner Telemetry Style')
plt.tight_layout()
plt.show()

driver_profiles_sorted['Driver_Style_Cluster'] = cluster_labels

```



By applying a second clustering step to these corner distributions, we classify drivers into distinct overall 'driving styles'.

1.9 PCA to show driver clusters

```
[34]: # Reduce driver profiles to 2D for the scatter plot
pca_drivers = PCA(n_components=2, random_state=42)
driver_features = driver_profiles.drop(columns=['Driver_Style_Cluster'])
pca_drivers_2d = pca_drivers.fit_transform(driver_features)

pca_df = pd.DataFrame(data=pca_drivers_2d, columns=['PC1', 'PC2'],
    ↪ index=driver_profiles.index)
pca_df['Driver_Style_Cluster'] = driver_profiles['Driver_Style_Cluster']

# Transform high-dimensional cluster centers into the 2D PCA space
centroids_2d = pca_drivers.transform(kmeans_drivers.cluster_centers_)

# Plot the results
plt.figure(figsize=(12, 8))
sns.set_style("whitegrid")
sns.scatterplot(x='PC1', y='PC2', hue='Driver_Style_Cluster', data=pca_df,
    ↪ palette='Set1', s=100)


# Plot the centroids as large black 'X' markers
plt.scatter(
    centroids_2d[:, 0],
    centroids_2d[:, 1],
    marker='X',
    s=300,
    color='black',
    label='Centroids',
    zorder=10
)

# Annotate each point with the driver number
for idx, row in pca_df.iterrows():
    plt.annotate(
        str(int(idx)),
        (row['PC1'], row['PC2']),
        textcoords="offset points",
        xytext=(5, 5),
        ha='left',
        fontsize=10,
        fontweight='bold'
    )

plt.title('Driver Style Clusters (2D PCA Projection)')
plt.xlabel('PC1 (Driver Profile Variance)')
plt.ylabel('PC2 (Driver Profile Variance)')
```

```
# Update legend to include the centroids
plt.legend(title='Driver Cluster & Centroids')
plt.tight_layout()
plt.show()
```

/Users/idomac/Documents/F1-Project/F1-Clustering-
 Project/venv/lib/python3.13/site-packages/sklearn/utils/validation.py:2827:
 UserWarning: X does not have valid feature names, but PCA was fitted with
 feature names
 warnings.warn(



The final visualization identifies three distinct groups of drivers, each representing a unique driving style.

Ultimately, using PCA beforehand led to more definitive clusters, effectively highlighting the distinct characteristics of each driving style.